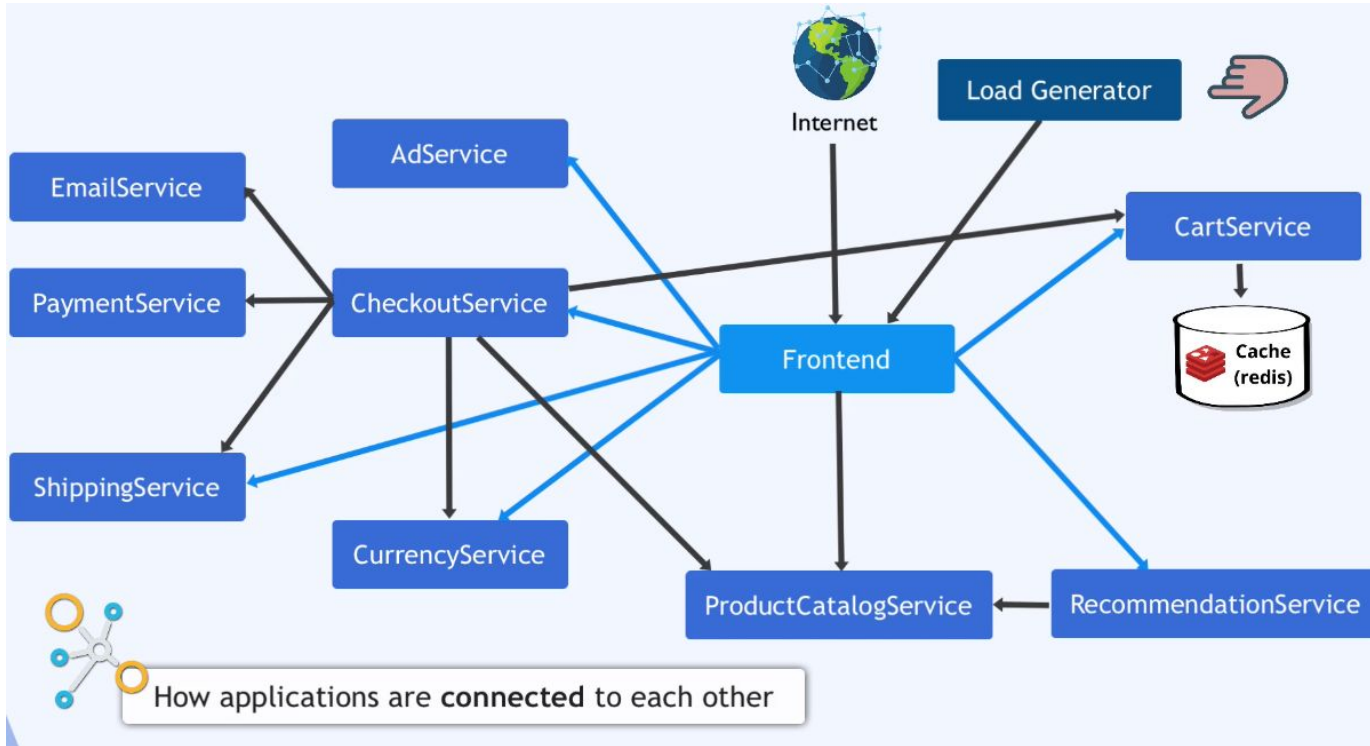


## Deploy microservice app into k8s cluster

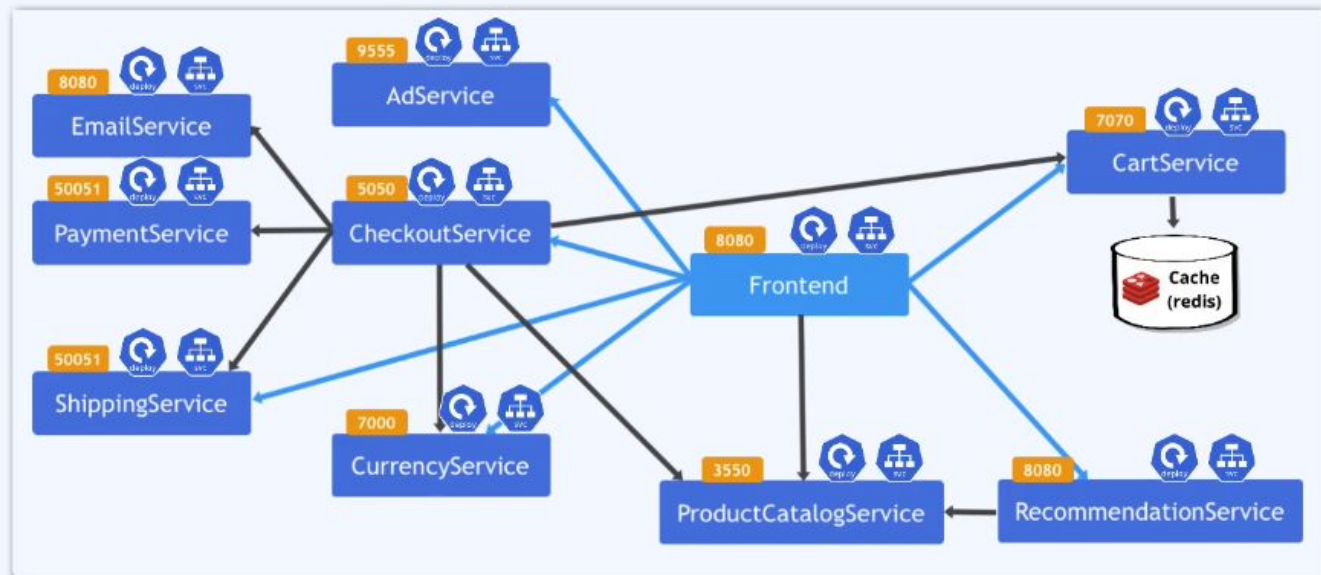


how applications are connected to each other  
image names for each microservice  
environment variables  
port  
deploy ms into single spaces



Create **Deployment Config**  
for **each Microservice**

Create **Service Config**  
for **each Microservice**



1. create a config file
  - a. skeleton
2. speed up k8s cluster in linode
  - a. download the kubeconfig file
3. set permission of kubeconfig file

```
[W]$ chmod 400 ~/Downloads/online-shop-kubeconfig.yaml  
[W]$ ls -l ~/Downloads/online-shop-kubeconfig.yaml  
-r-----@ 1 nanajanashia staff 2764 Jun 16 09:54 /Users  
config.yaml
```

#### 4. variable

```
[W]$ export KUBECONFIG=~/Downloads/online-shop-microservices-kubeconfig.yaml
```

5. kubectl get node check if the cluster right configured
6. kubectl create ns microservices
7. kubectl apply -f config.yaml -n microservices
8. kubectl get pod -n microservices
9. kubectl get svc -n microservices

linode

Linodes

Volumes

NodeBalancers

Firewalls

StackScripts

Images

Domains

Kubernetes

Object Storage

Longview

Marketplace

Account

Help & Support

Create

Search for Linodes, Volumes, NodeBalancers, Domains, Buckets, Tags...

A new version of Kubernetes is available (1.21).

Kubernetes / **online-shop-microservices**

Version 1.20

3 CPU Cores

Frankfurt, DE

\$30/month

6 GB RAM

150 GB Storage

Kubernetes API Endpoint:

https://7dc7925e-a17e-4b3f-8fbd-bb3724f8ac43.eu-central-2.linode.lke.net:443

Kubeconfig:

online-shop-microservices-kubeconfig.yaml

Download

Copy

Node Pools

Linode 2 GB

Linode ^	Status ^	IP Address ^
<div></div> lke28633-42767-60c34bde69f3	Running	193.195.240.188
<div></div> lke28633-42767-60c34bdecc8a	Running	194.233.164.143
<div></div> lke28633-42767-60c34bdf2dc8	Running	194.195.241.114

Not Secure | 194.195.240.188:30007

Free shipping with \$75 purchase!



Recommendation Service needs to know the endpoint of ProductCatalogService

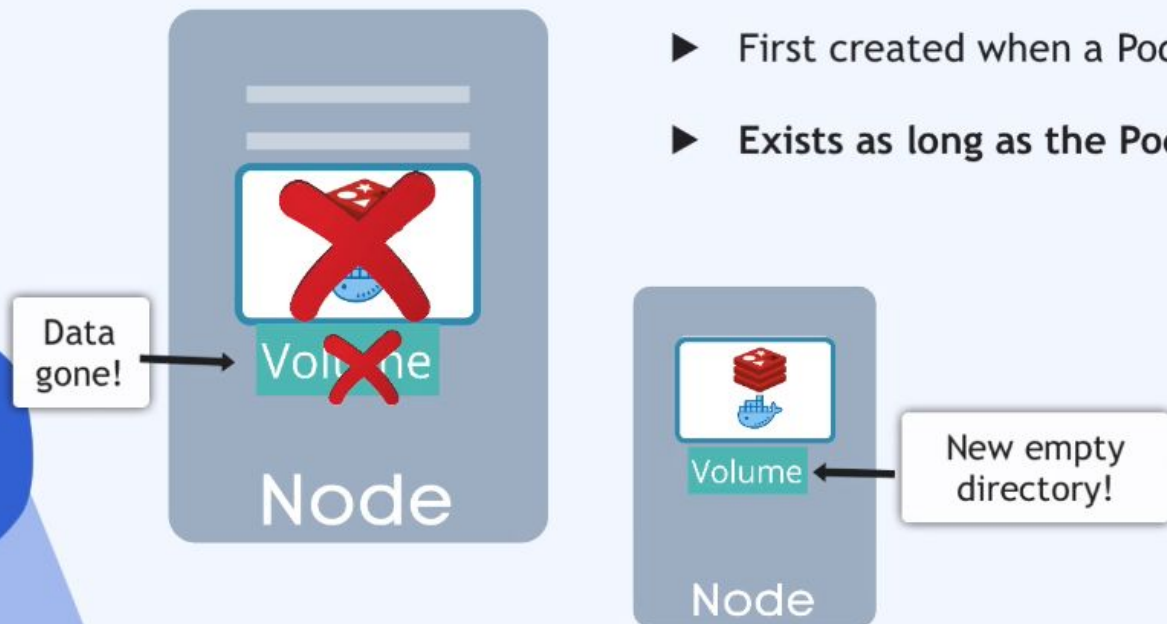
Endpoint = Service Name + Service Port





**emptyDir**

- ▶ Is initially empty
- ▶ First created when a Pod is assigned to a Node
- ▶ Exists as long as the Pod is running







## emptyDir

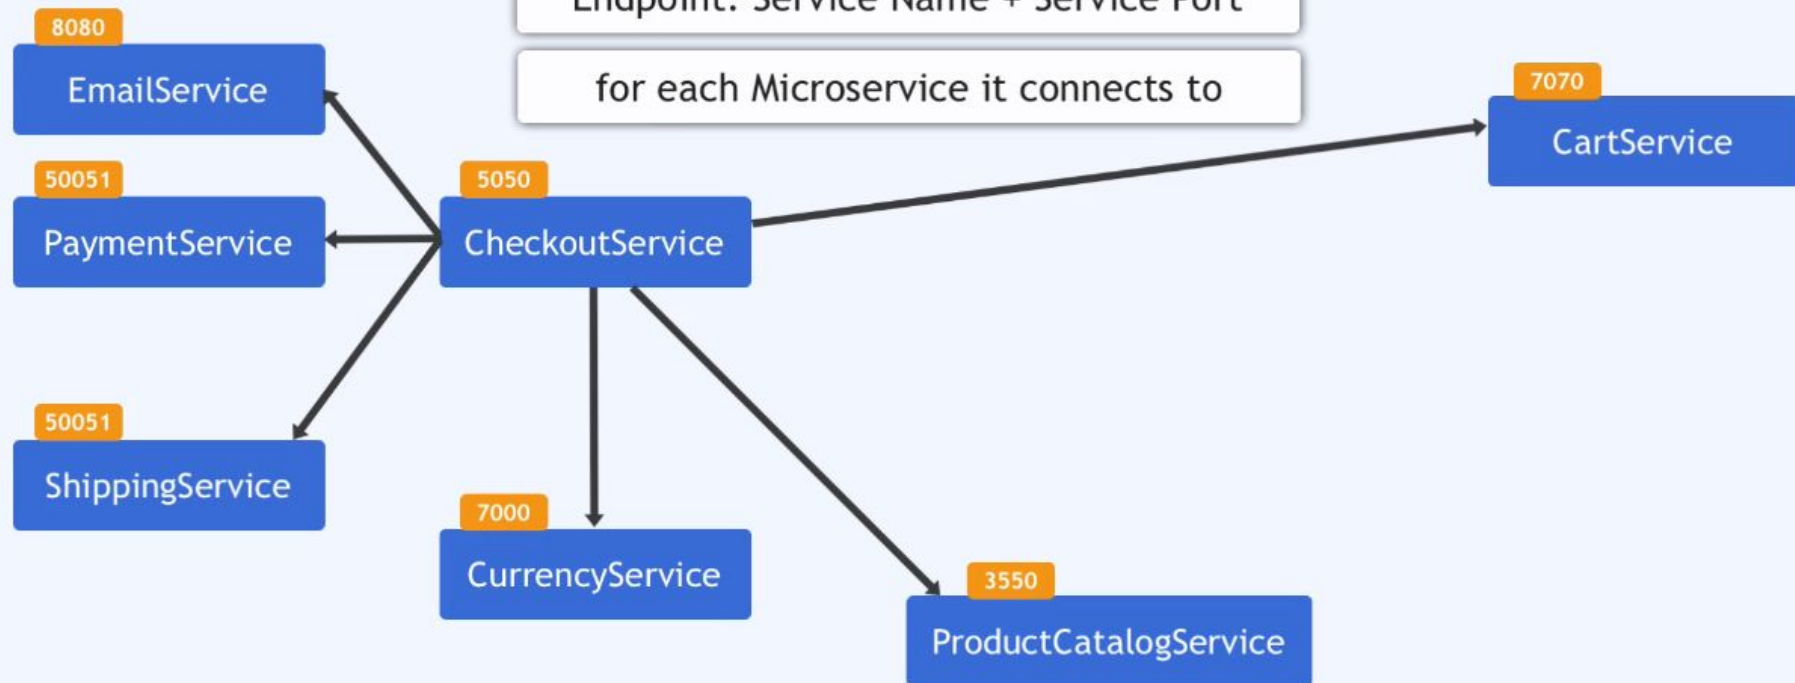
- ▶ Is initially empty
- ▶ First created when a Pod is assigned to a Node
- ▶ **Exists as long as the Pod is running**
- ▶ Container crashing does NOT remove a Pod from a Node
- ▶ Therefore, **data safe across container crashes!**



CheckoutService needs to know:

Endpoint: Service Name + Service Port

for each Microservice it connects to



```
88         port = os.Getenv("PORT")
89     }
90
91     svc := new(checkoutService)
92     mustMapEnv(&svc.shippingSvcAddr, "SHIPPING_SERVICE_ADDR")
93     mustMapEnv(&svc.productCatalogSvcAddr, "PRODUCT_CATALOG_SERVICE_ADDR")
94     mustMapEnv(&svc.cartSvcAddr, "CART_SERVICE_ADDR")
95     mustMapEnv(&svc.currencySvcAddr, "CURRENCY_SERVICE_ADDR")
96     mustMapEnv(&svc.emailSvcAddr, "EMAIL_SERVICE_ADDR")
97     mustMapEnv(&svc.paymentSvcAddr, "PAYMENT_SERVICE_ADDR")
98
```

```
348     image: gcr.io/google-samples/microservices-demo/checkoutservice
349     ports:
350     - containerPort: 5050
351     env:
352     - name: PORT
353       value: "5050"
354     - name: PRODUCT_CATALOG_SERVICE_ADDR
355     ---
356     apiVersion: v1
357     kind: Service
358     metadata:
359       name: checkoutservice
360     spec:
361       type: ClusterIP
```


```
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: emailservice
spec:
  selector:
    matchLabels:
      app: emailservice
  template:
    metadata:
      labels:
        app: emailservice
    spec:
      containers:
        - name: service
          image: gcr.io/google-samples/emailservice
          ports:
            - containerPort: 8080
          env:
            - name: PORT
              value: "8080"
```

```
}
---
apiVersion: v1
kind: Service
metadata:
  name: emailservice
spec:
  type: ClusterIP
  selector:
    app: emailservice
  ports:
    - protocol: TCP
      port: 5000
      targetPort: 8080
```

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: cartservice
spec:
  selector:
    matchLabels:
      app: cartservice
  template:
    metadata:
      labels:
        app: cartservice
    spec:
      containers:
        - name: service
          image: gcr.io/google-samples/cartservice
          ports:
            - containerPort: 7070
          env:
            - name: PORT
              value: "7070"
            - name: REDIS_ADDR
              value: "redis-cart:6379"

```



```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis-cart
spec:
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: redis
          image: redis:alpine
          ports:
            - containerPort: 6379
          volumeMounts:
            - mountPath: /data
              name: redis-data
      volumes:
        - name: redis-data
          emptyDir: {}

```

```


apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  selector:
    matchLabels:
      app: frontend
  template:
    metadata:
      labels:
        app: frontend
    spec:
      containers:
        - name: service
          image: gcr.io/google-samples/frontend
          ports:
            - containerPort: 8080
          env:
            - name: PORT
              value: "8080"
            - name: SHIPPING_SERVICE_ADDR
              value: "shippingservice:50051"
            - name: PRODUCT_CATALOG_SERVICE_ADDR
              value: "productcatalogservice:3550"
            - name: CART_SERVICE_ADDR

```

```

---
apiVersion: v1
kind: Service
metadata:
  name: frontend
spec:
  type: NodePort
  selector:
    app: frontend
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 8080
      nodePort: 30007
---

```



best

practices

1. define image version
2. liveness
3. readiness



when the container is not running  
kubectl logs checkout  
disable functions

```
config.yaml x ! tmp.yaml
! config.yaml > {} spec > {} template > {} spec > [ ] containers > {} 0 > [ ] env
66 - name: service
67   image: gcr.io/google-samples/microservices-demo/recommendati
68   ports:
69     - containerPort: 8080
70   env:
71     - name: PORT
72       value: "8080"
73     - name: PRODUCT_CATALOG_SERVICE_ADDR
74       value: "productcatalogservice:3550"
75     - name: DISABLE_TRACING
76       value: "1"
77     - name: DISABLE_PROFILER
78       value: "1"
79     - name: DISABLE_DEBUGGER
80       value: "1"
81   readinessProbe:
```

```
[W]$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
emailservice-5787b7575-jhtrk	1/1	Running	0	36s
productcatalogservice-54d7fc98bb-nj7f2	1/1	Running	0	36s
recommendationservice-68956c9949-2z5md	1/1	Running	0	36s

```
! config.yaml > {} spec > {} template > {} spec > [ ] containers > {} 0 > {} readinessProbe > {} exec > [ ] command
```

## 2) TCP probes

Kubelet makes probe connection at the node, not in the pod

```
23         value: "8080"
```

```
26     - name: DISABLE_P
```

```
27         value: "1"
```

```
28     readinessProbe:
```

```
29         periodSeconds: 5
```

```
30         exec:
```

```
31             command: ["/bin/grpc_health_probe", "-addr=:8080"]
```

```
32     livenessProbe:
```

```
33         periodSeconds: 5
```

```
34         exec:
```

```
35             command: ["/bin/grpc_health_probe", "-addr=:8080"]
```

```
36     ---
```

! config.yaml X

! config.yaml > {} spec > {} template > {} spec > [ ] containers > {} 0 > {} livenessProbe > {} tcpSocket > #

```
360   template:
361     metadata:
362       labels:
363         app: redis-cart
364     spec:
365       containers:
366       - name: redis
367         image: redis:alpine
368         ports:
369         - containerPort: 6379
370         readinessProbe:
371           periodSeconds: 5
372           tcpSocket:
373             port: 6379
374         livenessProbe:
375           periodSeconds: 5
376           tcpSocket:
377             port: 6379
378         volumeMounts:
379         - name: redis-data
```

methods:  
command  
tcpsocket  
http

```
spec:
  containers:
  - name: redis
    image: redis:alpine
    ports:
    - containerPort: 6379
    readinessProbe:
      initialDelaySeconds: 5
      periodSeconds: 5
      tcpSocket:
        port: 6379
    livenessProbe:
      initialDelaySeconds: 5
      periodSeconds: 5
      tcpSocket:
```

wait 5 seconds before performing probe

### 3) HTTP probes

Kubelet sends an HTTP request to specified path and port

```
! config.yaml > {} spec > {} template > {} spec > [ ] containers > {} 0 > {} livenessProbe > {} tcpSocket > # port
118 containers:
119   - name: service
122     - containerPort: 3550
123     env:
124       - name: PORT
125         value: "3550"
126     livenessProbe:
127       periodSeconds: 5
128       exec:
129         command: ["/bin/grpc_health_probe", "-addr=:3550"]
130     readinessProbe:
131       periodSeconds: 5
```

```
yaml > {} spec > {} template > {} spec > [ ] containers > {} 0 > {} livene
  app: productcatalogservice
spec:
  containers:
  - name: service
    image: gcr.io/google-samples/microservice
    ports:
    - containerPort: 3550
    env:
    - name: PORT
      value: "3550"
    livenessProbe:
      periodSeconds: 5
      httpGet:
        path: /health
        port: 3550
    readinessProbe:
      periodSeconds: 5
      exec:
        command: ["/bin/grpc_health_probe", "
```

define resource

```

- name: PORT
  value: "8080"
livenessProbe:
  periodSeconds: 5
  exec:
    command: ["/bin/grpc_health_probe", "-ad
readinessProbe:
  periodSeconds: 5
  exec:
    command: ["/bin/grpc_health_probe", "-ad
resources:
  requests:
    cpu: 100m
    memory: 64Mi

```

```

  command: ["/bin/grpc_heal
resources:
  requests:
    cpu: 100m
    memory: 64Mi
  limits:
    cpu: 200m
    memory: 128Mi

```



CPU resources are  
defined in **millicores**

Memory resources are  
defined in **bytes**

```

Version: v1
d: Service
adata:
ame: emailservice

```

m = "millicore"

Mi = "mebibyte"



# Why Resource Limits are important?

What if application needs more resources?



Needing more..




Request 1 CPU

- ▶ Container will consume more than the requested resources
- ▶ If not limited, container could consume all the Node's resources

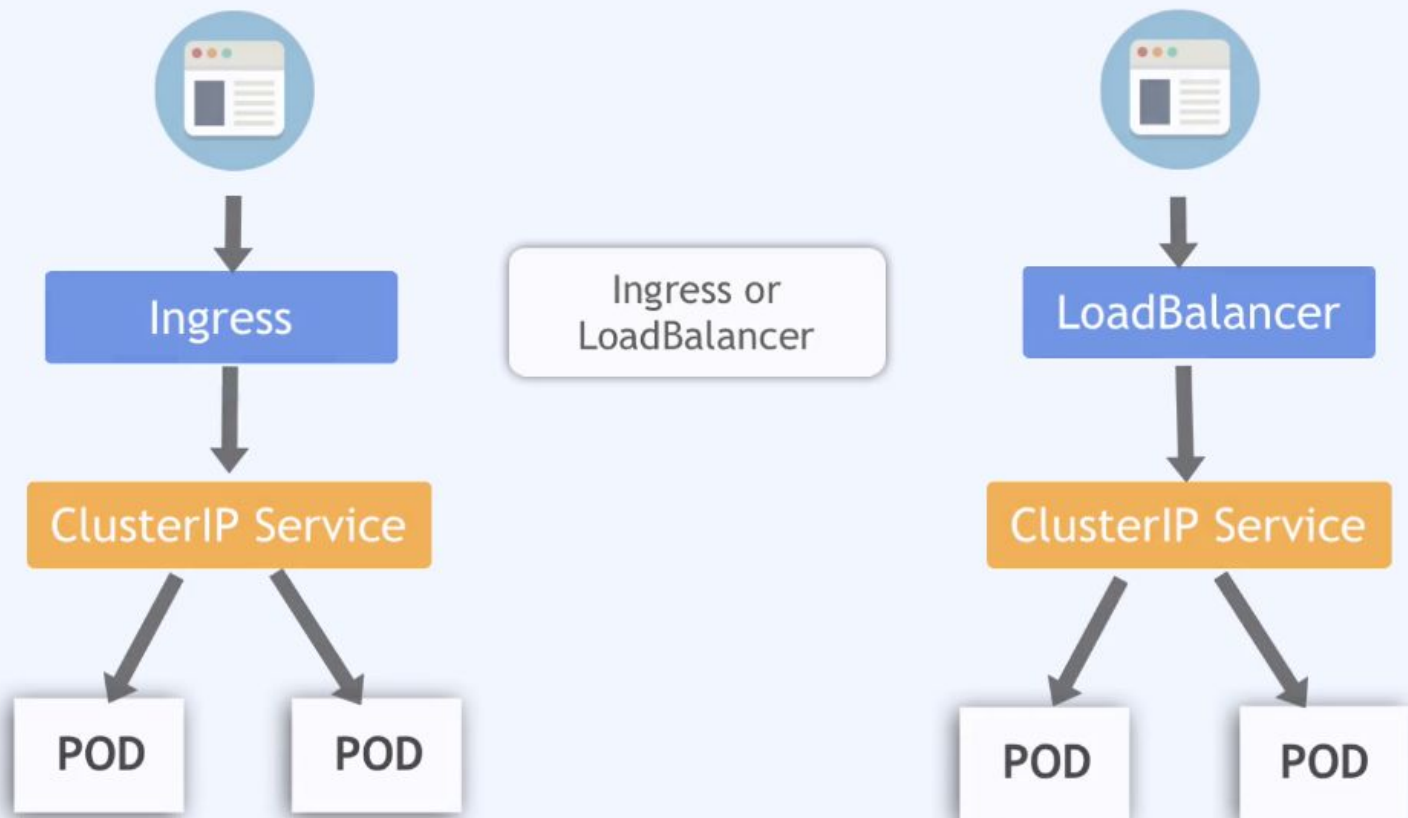


# Why it's a Bad Practice?



```
apiVersion: v1
kind: Service
metadata:
  name: frontend-external
spec:
  type: LoadBalancer
  selector:
    app: frontend
  ports:
    - name: http
      port: 80
      targetPort: 8080
```

# Ingress as alternative



# Why it's important?

- Custom Identifier for your components

## 1. Group Pods with Labels

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: emailservice
spec:
  selector:
    matchLabels:
      app: emailservice
  template:
    metadata:
      labels:
        app: emailservice
    spec: ...
```

## 2. Reference in Service Component

```
apiVersion: v1
kind: Service
metadata:
  name: emailservice
spec:
  type: ClusterIP
  selector:
    app: emailservice
  ports:
    - protocol: TCP
      port: 5000
      targetPort: 8080
```





